

Random Forest Project

For this project we will be exploring publicly available data from [LendingClub.com](https://www.lendingclub.com) (www.lendingclub.com). Lending Club connects people who need money (borrowers) with people who have money (investors). Hopefully, as an investor you would want to invest in people who showed a profile of having a high probability of paying you back. We will try to create a model that will help predict this.

Lending club had a [very interesting year in 2016](https://en.wikipedia.org/wiki/Lending_Club#2016) (https://en.wikipedia.org/wiki/Lending_Club#2016), so let's check out some of their data and keep the context in mind. This data is from before they even went public.

We will use lending data from 2007-2010 and be trying to classify and predict whether or not the borrower paid back their loan in full. You can download the data from [here](https://www.lendingclub.com/info/download-data.action) (<https://www.lendingclub.com/info/download-data.action>) or just use the csv already provided. It's recommended you use the csv provided as it has been cleaned of NA values.

Here are what the columns represent:

- credit.policy: 1 if the customer meets the credit underwriting criteria of LendingClub.com, and 0 otherwise.
- purpose: The purpose of the loan (takes values "credit_card", "debt_consolidation", "educational", "major_purchase", "small_business", and "all_other").
- int.rate: The interest rate of the loan, as a proportion (a rate of 11% would be stored as 0.11). Borrowers judged by LendingClub.com to be more risky are assigned higher interest rates.
- installment: The monthly installments owed by the borrower if the loan is funded.
- log.annual.inc: The natural log of the self-reported annual income of the borrower.
- dti: The debt-to-income ratio of the borrower (amount of debt divided by annual income).
- fico: The FICO credit score of the borrower.
- days.with.cr.line: The number of days the borrower has had a credit line.
- revol.bal: The borrower's revolving balance (amount unpaid at the end of the credit card billing cycle).
- revol.util: The borrower's revolving line utilization rate (the amount of the credit line used relative to total credit available).
- inq.last.6mths: The borrower's number of inquiries by creditors in the last 6 months.
- delinq.2yrs: The number of times the borrower had been 30+ days past due on a payment in the past 2 years.
- pub.rec: The borrower's number of derogatory public records (bankruptcy filings, tax liens, or judgments).

Import Libraries

Import the usual libraries for pandas and plotting. You can import sklearn later on.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Get the Data

**** Use pandas to read loan_data.csv as a dataframe called loans.****

```
In [2]: loans = pd.read_csv('loan_data.csv')
```

**** Check out the info(), head(), and describe() methods on loans.****

```
In [3]: loans.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   credit.policy          9578 non-null  int64  
1   purpose                9578 non-null  object  
2   int.rate               9578 non-null  float64 
3   installment            9578 non-null  float64 
4   log.annual.inc         9578 non-null  float64 
5   dti                   9578 non-null  float64 
6   fico                  9578 non-null  int64  
7   days.with.cr.line      9578 non-null  float64 
8   revol.bal              9578 non-null  int64  
9   revol.util             9578 non-null  float64 
10  inq.last.6mths         9578 non-null  int64  
11  delinq.2yrs            9578 non-null  int64  
12  pub.rec                9578 non-null  int64  
13  not.fully.paid         9578 non-null  int64  
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```

In [4]: `loans.describe()`

Out[4]:

	policy	int.rate	installment	log.annual.inc	dti	fico	days.with.cr.line	revol.
count	1000	9578.000000	9578.000000	9578.000000	9578.000000	9578.000000	9578.000000	9.578000e
mean	4970	0.122640	319.089413	10.932117	12.606679	710.846314	4560.767197	1.691396e
std	3245	0.026847	207.071301	0.614813	6.883970	37.970537	2496.930377	3.375619e
min	1000	0.060000	15.670000	7.547502	0.000000	612.000000	178.958333	0.000000e
25%	1000	0.103900	163.770000	10.558414	7.212500	682.000000	2820.000000	3.187000e
50%	1000	0.122100	268.950000	10.928884	12.665000	707.000000	4139.958333	8.596000e
75%	1000	0.140700	432.762500	11.291293	17.950000	737.000000	5730.000000	1.824950e
max	1000	0.216400	940.140000	14.528354	29.960000	827.000000	17639.958330	1.207359e

In [5]: `loans.head()`

Out[5]:

	policy	purpose	int.rate	installment	log.annual.inc	dti	fico	days.with.cr.line	revol.bal
0	1	debt_consolidation	0.1189	829.10	11.350407	19.48	737	5639.958333	28854
1	1	credit_card	0.1071	228.22	11.082143	14.29	707	2760.000000	33623
2	1	debt_consolidation	0.1357	366.86	10.373491	11.63	682	4710.000000	3511
3	1	debt_consolidation	0.1008	162.34	11.350407	8.10	712	2699.958333	33667
4	1	credit_card	0.1426	102.92	11.299732	14.97	667	4066.000000	4740


Exploratory Data Analysis

Let's do some data visualization! We'll use seaborn and pandas built-in plotting capabilities, but feel free to use whatever library you want. Don't worry about the colors matching, just worry about getting the main idea of the plot.

**** Create a histogram of two FICO distributions on top of each other, one for each credit.policy outcome.****


Note: This is pretty tricky, feel free to reference the solutions. You'll probably need one line of code for each histogram, I also recommend just using pandas built in .hist()

```
In [6]: plt.figure(figsize=[10,6])
sns.set_style('darkgrid')
sns.set_context('paper')
loans[loans['credit.policy']==1]['fico'].hist(bins=30,alpha=0.5,color='b',label='Credit.Policy=1')
loans[loans['credit.policy']==0]['fico'].hist(bins=30,alpha=0.5,color='r',label='Credit.Policy=0')
plt.legend();
plt.xlabel('FICO');
```



**** Create a similar figure, except this time select by the not.fully.paid column.****


```
In [7]: plt.figure(figsize=[10,6])
sns.set_style('darkgrid')
sns.set_context('paper')
loans[loans['not.fully.paid']==1]['fico'].hist(bins=30,alpha=0.5,color='b',label='not.fully.paid=1')
loans[loans['not.fully.paid']==0]['fico'].hist(bins=30,alpha=0.5,color='r',label='not.fully.paid=0')
plt.legend();
plt.xlabel('FICO');
```



**** Create a countplot using seaborn showing the counts of loans by purpose, with the color hue defined by not.fully.paid. ****

```
In [8]: plt.figure(figsize=[11,7])  
sns.countplot(x='purpose',hue='not.fully.paid',data=loans,palette='Set1')
```


Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x1dfc0b70108>



**** Let's see the trend between FICO score and interest rate. Recreate the following jointplot.****

```
In [9]: sns.jointplot(x='fico', y='int.rate',data=loans)
```

```
Out[9]: <seaborn.axisgrid.JointGrid at 0x1dfc0b3a508>
```



**** Create the following Implots to see if the trend differed between not.fully.paid and credit.policy.
Check the documentation for Implot() if you can't figure out how to separate it into columns.****

```
In [19]: sns.lmplot(x="fico", y="int.rate", hue= 'credit.policy', col="not.fully.paid", data=loans)
```

```
Out[19]: <seaborn.axisgrid.FacetGrid at 0x1dfc2ddd648>
```



Setting up the Data

Let's get ready to set up our data for our Random Forest Classification Model!

Check loans.info() again.

```
In [20]: loans.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   credit.policy          9578 non-null   int64
1   purpose               9578 non-null   object
2   int.rate              9578 non-null   float64
3   installment           9578 non-null   float64
4   log.annual.inc        9578 non-null   float64
5   dti                   9578 non-null   float64
6   fico                  9578 non-null   int64
7   days.with.cr.line     9578 non-null   float64
8   revol.bal             9578 non-null   int64
9   revol.util            9578 non-null   float64
10  inq.last.6mths        9578 non-null   int64
11  delinq.2yrs           9578 non-null   int64
12  pub.rec               9578 non-null   int64
13  not.fully.paid        9578 non-null   int64
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```


Categorical Features

Notice that the **purpose** column is categorical

That means we need to transform them using dummy variables so sklearn will be able to understand them. Let's do this in one clean step using `pd.get_dummies`.

Let's show you a way of dealing with these columns that can be expanded to multiple categorical features if necessary.

Create a list of 1 element containing the string 'purpose'. Call this list `cat_feats`.

```
In [21]: cat_feats= ['purpose']
```

Now use `pd.get_dummies(loans,columns=cat_feats,drop_first=True)` to create a fixed larger dataframe that has new feature columns with dummy variables. Set this dataframe as `final_data`.

```
In [25]: final_data=pd.get_dummies(loans,columns=cat_feats,drop_first=True)
```

```
In [26]: final_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 19 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   credit.policy                        9578 non-null   int64
 1   int.rate                            9578 non-null   float64
 2   installment                         9578 non-null   float64
 3   log.annual.inc                     9578 non-null   float64
 4   dti                                 9578 non-null   float64
 5   fico                                9578 non-null   int64
 6   days.with.cr.line                  9578 non-null   float64
 7   revol.bal                          9578 non-null   int64
 8   revol.util                         9578 non-null   float64
 9   inq.last.6mths                     9578 non-null   int64
10   delinq.2yrs                        9578 non-null   int64
11   pub.rec                            9578 non-null   int64
12   not.fully.paid                     9578 non-null   int64
13   purpose_credit_card                 9578 non-null   uint8
14   purpose_debt_consolidation          9578 non-null   uint8
15   purpose_educational                 9578 non-null   uint8
16   purpose_home_improvement            9578 non-null   uint8
17   purpose_major_purchase              9578 non-null   uint8
18   purpose_small_business              9578 non-null   uint8
dtypes: float64(6), int64(7), uint8(6)
memory usage: 1.0 MB
```

Train Test Split

Now its time to split our data into a training set and a testing set!

**** Use sklearn to split your data into a training set and a testing set as we've done in the past.****

```
In [27]: from sklearn.model_selection import train_test_split
```

```
In [32]: X = final_data.drop('not.fully.paid',axis=1)
y=final_data['not.fully.paid']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_
```

Training a Decision Tree Model

Let's start by training a single decision tree first!

**** Import DecisionTreeClassifier****

```
In [33]: from sklearn.tree import DecisionTreeClassifier
```

Create an instance of DecisionTreeClassifier() called dtree and fit it to the training data.

```
In [34]: dtree = DecisionTreeClassifier()
```

```
In [35]: dtree.fit(X_train,y_train)
```

```
Out[35]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                                max_depth=None, max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort='deprecated',
                                random_state=None, splitter='best')
```

Predictions and Evaluation of Decision Tree

Create predictions from the test set and create a classification report and a confusion matrix.

```
In [36]: pred = dtree.predict(X_test)
```

```
In [37]: from sklearn.metrics import confusion_matrix,classification_report
```

```
In [38]: print(classification_report(pred,y_test))
```

	precision	recall	f1-score	support
0	0.84	0.85	0.85	2620
1	0.24	0.22	0.23	541
accuracy			0.74	3161
macro avg	0.54	0.54	0.54	3161
weighted avg	0.74	0.74	0.74	3161

```
In [39]: print(confusion_matrix(pred,y_test))
```

```
[[2230  390]
 [ 420  121]]
```

Training the Random Forest model

Now its time to train our model!

Create an instance of the RandomForestClassifier class and fit it to our training data from the previous step.

```
In [40]: from sklearn.ensemble import RandomForestClassifier
```

```
In [41]: rfor = RandomForestClassifier(n_estimators=100)
```

```
In [42]: rfor.fit(X_train,y_train)
```

```
Out[42]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                                criterion='gini', max_depth=None, max_features='auto',
                                max_leaf_nodes=None, max_samples=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100,
                                n_jobs=None, oob_score=False, random_state=None,
                                verbose=0, warm_start=False)
```

Predictions and Evaluation

Let's predict off the y_test values and evaluate our model.

**** Predict the class of not.fully.paid for the X_test data.****

```
In [43]: pred_R = rfor.predict(X_test)
```

Now create a classification report from the results. Do you get anything strange or some sort of warning?

```
In [44]: print(classification_report(pred_R,y_test))
```

	precision	recall	f1-score	support
0	0.99	0.84	0.91	3136
1	0.02	0.32	0.03	25
accuracy			0.84	3161
macro avg	0.50	0.58	0.47	3161
weighted avg	0.99	0.84	0.90	3161

Show the Confusion Matrix for the predictions.

```
In [45]: print(confusion_matrix(pred_R,y_test))
```

```
[[2633  503]
 [  17    8]]
```

```
In [36]:
```

Great Job!